

```
In [ ]: import os
        #os.getcwd()
        import warnings
        warnings.filterwarnings('ignore')

        import sys
        import numpy as np
        import pandas as pd
        import matplotlib.pyplot as plt
        %matplotlib inline
        import seaborn as sns
        sns.set()
        import sklearn
        import scipy.integrate as integrate
        import scipy

        import statsmodels.api as sm
        from statsmodels.formula.api import ols
        from sklearn.preprocessing import KBinsDiscretizer
        from sklearn.feature_selection import SelectKBest, chi2
        from sklearn.model_selection import train_test_split
        from sklearn.preprocessing import StandardScaler
        from sklearn.preprocessing import MinMaxScaler
        from sklearn.feature_selection import RFE
        from sklearn.feature_selection import RFECV

        # importing algorithm, evaluation, and model Libraries
        from sklearn.feature_extraction.text import CountVectorizer
        from sklearn.model_selection import train_test_split
        from sklearn.linear_model import LogisticRegression
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.ensemble import GradientBoostingClassifier
        from xgboost import XGBClassifier
        from sklearn.linear_model import RidgeClassifier
        from sklearn.model_selection import GridSearchCV
        from sklearn.svm import SVC

        import sklearn.metrics as metrics
        from sklearn.metrics import accuracy_score
        from sklearn.metrics import confusion_matrix
        from sklearn.metrics import recall_score
        from sklearn.model_selection import cross_val_score
        from statsmodels.stats.outliers_influence import variance_inflation_factor
```

```
In [ ]: churn_data = pd.read_csv('churn.csv')
```

```
In [ ]: churn_data.head()
```

Out []:

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	CallService	MultipleConn
--	------------	--------	---------------	---------	------------	--------	-------------	--------------

0	2907-ILJBN	Female	0.0	Yes	Yes	11.0	Yes	
1	3896-RCYYE	Female	0.0	No	No	67.0	No	No phone
2	9764-REAFF	Female	0.0	Yes	No	59.0	Yes	
3	6651-RLGGM	Male	0.0	Yes	Yes	67.0	Yes	
4	5879-SESNB	Female	0.0	No	No	11.0	Yes	

5 rows × 21 columns

In []: churn_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12335 entries, 0 to 12334
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   customerID                           12335 non-null  object
1   gender                               12335 non-null  object
2   SeniorCitizen                         12335 non-null  float64
3   Partner                               12335 non-null  object
4   Dependents                           12335 non-null  object
5   tenure                               12335 non-null  float64
6   CallService                           12335 non-null  object
7   MultipleConnections                   12335 non-null  object
8   InternetConnection                   12335 non-null  object
9   OnlineSecurity                        12335 non-null  object
10  OnlineBackup                          12335 non-null  object
11  DeviceProtectionService               12335 non-null  object
12  TechnicalHelp                         12335 non-null  object
13  OnlineTV                              12335 non-null  object
14  OnlineMovies                          12335 non-null  object
15  Agreement                             12335 non-null  object
16  BillingMethod                          12335 non-null  object
17  PaymentMethod                         12335 non-null  object
18  MonthlyServiceCharges                 12335 non-null  float64
19  TotalAmount                           12335 non-null  float64
20  Churn                                 12335 non-null  object
dtypes: float64(4), object(17)
memory usage: 2.0+ MB
```

In []: churn_data = churn_data.drop_duplicates()
churn_data.shape

Out []: (9328, 21)

data has no null values

```
In [ ]: churn_data.isnull().sum()
```

```
Out[ ]: customerID          0
gender                    0
SeniorCitizen            0
Partner                  0
Dependents                0
tenure                   0
CallService              0
MultipleConnections      0
InternetConnection       0
OnlineSecurity           0
OnlineBackup             0
DeviceProtectionService  0
TechnicalHelp            0
OnlineTV                 0
OnlineMovies             0
Agreement                0
BillingMethod            0
PaymentMethod            0
MonthlyServiceCharges    0
TotalAmount              0
Churn                    0
dtype: int64
```

dropping customerID

```
In [ ]: churn_data = churn_data.drop(['customerID'],axis=1)
churn_data.isnull().sum()
```

```
Out[ ]: gender          0
SeniorCitizen          0
Partner                0
Dependents              0
tenure                 0
CallService            0
MultipleConnections     0
InternetConnection     0
OnlineSecurity          0
OnlineBackup           0
DeviceProtectionService 0
TechnicalHelp          0
OnlineTV               0
OnlineMovies           0
Agreement              0
BillingMethod          0
PaymentMethod          0
MonthlyServiceCharges  0
TotalAmount            0
Churn                  0
dtype: int64
```

```
In [ ]: y=churn_data[['Churn']]
        x=churn_data.drop(['Churn'],axis=1)
        y.shape
```

```
Out[ ]: (9328, 1)
```

finding the % of churn-Yes Or the mean of Y = Yes

```
In [ ]: y['Churn'].value_counts()
```

```
Out[ ]: Yes    5607
        No     3721
        Name: Churn, dtype: int64
```

```
In [ ]: # finding genders distribution with churners
        churn_percent = (y['Churn'].value_counts() / len(y) * 100)
        print(churn_percent)
        y_temp = pd.DataFrame(y)
        y_temp = np.where(y_temp['Churn']=='No',0,1)
        print('churn_percent ',y_temp.mean())
```

```
Yes    60.109348
No     39.890652
Name: Churn, dtype: float64
churn_percent 0.6010934819897084
```

spitting the features categroical & numeric separtely

```
In [ ]: x_num=x.select_dtypes(include="number")
        x_cat=x.select_dtypes(include="object")
```

numeric fields

```
In [ ]: x_num.columns
```

```
Out[ ]: Index(['SeniorCitizen', 'tenure', 'MonthlyServiceCharges', 'TotalAmount'], dtype='object')
```

```
In [ ]: x_num.describe()
```

Out []:

	SeniorCitizen	tenure	MonthlyServiceCharges	TotalAmount
count	9328.000000	9328.000000	9328.000000	9328.000000
mean	0.163885	26.642965	61.341895	1839.159069
std	0.351640	22.491378	27.661949	2008.858997
min	0.000000	1.000000	18.250000	18.800000
25%	0.000000	6.700780	35.392830	288.599227
50%	0.000000	20.325530	64.220312	1052.258266
75%	0.000000	44.345050	84.070699	2777.241953
max	1.000000	72.000000	118.750000	8684.800000

In []: `x_num['SeniorCitizen'].value_counts()`

Out []: 0.000000 7420
 1.000000 1155
 0.660002 1
 0.208768 1
 0.437521 1
 ...
 0.087772 1
 0.353116 1
 0.387478 1
 0.031710 1
 0.433856 1
 Name: SeniorCitizen, Length: 755, dtype: int64

Senior citizen should be Yes or No (1 or 0). But we have some float values here. Incorrect records are about 753 rows. Close to 6% of the data. So replacing them with mode, which is 0.

In []: `x_num['SeniorCitizen'] = np.where((x_num['SeniorCitizen'] > 0) & (x_num['SeniorC
 x['SeniorCitizen'].value_counts()`

Out []: 0.000000 7420
 1.000000 1155
 0.660002 1
 0.208768 1
 0.437521 1
 ...
 0.087772 1
 0.353116 1
 0.387478 1
 0.031710 1
 0.433856 1
 Name: SeniorCitizen, Length: 755, dtype: int64

In []: `x_cat.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 9328 entries, 0 to 12334
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   gender                                9328 non-null   object
1   Partner                               9328 non-null   object
2   Dependents                            9328 non-null   object
3   CallService                           9328 non-null   object
4   MultipleConnections                   9328 non-null   object
5   InternetConnection                    9328 non-null   object
6   OnlineSecurity                         9328 non-null   object
7   OnlineBackup                          9328 non-null   object
8   DeviceProtectionService               9328 non-null   object
9   TechnicalHelp                         9328 non-null   object
10  OnlineTV                              9328 non-null   object
11  OnlineMovies                          9328 non-null   object
12  Agreement                             9328 non-null   object
13  BillingMethod                          9328 non-null   object
14  PaymentMethod                         9328 non-null   object
dtypes: object(15)
memory usage: 1.1+ MB
```

```
In [ ]: x_num.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 9328 entries, 0 to 12334
Data columns (total 4 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   SeniorCitizen                        9328 non-null   float64
1   tenure                               9328 non-null   float64
2   MonthlyServiceCharges                9328 non-null   float64
3   TotalAmount                          9328 non-null   float64
dtypes: float64(4)
memory usage: 364.4 KB
```

treating outliers for numeric columns tenure, MonthlyServiceCharges, TotalAmount

```
In [ ]: upper_limit = x_num['MonthlyServiceCharges'].quantile(0.99)
lower_limit = x_num['MonthlyServiceCharges'].quantile(0.01)
x_num['MonthlyServiceCharges'] = np.where(x_num['MonthlyServiceCharges'] >= upper_limit, lower_limit, x_num['MonthlyServiceCharges'])
```

```
In [ ]: upper_limit = x_num['tenure'].quantile(0.99)
lower_limit = x_num['tenure'].quantile(0.01)
x_num['tenure'] = np.where(x_num['tenure'] >= upper_limit, lower_limit, x_num['tenure'])
```

```
In [ ]: upper_limit = x_num['TotalAmount'].quantile(0.99)
lower_limit = x_num['TotalAmount'].quantile(0.01)
x_num['TotalAmount'] = np.where(x_num['TotalAmount'] >= upper_limit, lower_limit, x_num['TotalAmount'])
```

```
In [ ]: # checking after outlier treatment
x_num.describe()
```

Out[]:

	SeniorCitizen	tenure	MonthlyServiceCharges	TotalAmount
count	9328.000000	9328.000000	9328.000000	9328.000000
mean	0.123821	26.642965	61.317044	1835.677311
std	0.329395	22.491378	27.607443	1998.063713
min	0.000000	1.000000	19.300000	20.200000
25%	0.000000	6.700780	35.392830	288.599227
50%	0.000000	20.325530	64.220312	1052.258266
75%	0.000000	44.345050	84.070699	2777.241953
max	1.000000	72.000000	112.236500	7780.474000

discretizing tenure column


```
In [ ]: x_tenure_binned = x_num['tenure']
x_tenure_binned = x_tenure_binned.values.reshape(-1, 1)
est = KBinsDiscretizer(n_bins=6, encode='ordinal', strategy='uniform')
Xt = est.fit_transform(x_tenure_binned)
```

```
In [ ]: x_num['tenure'] = Xt
x_num['tenure'].shape
```

Out[]: (9328,)

data visualization of numeric to numeric and selecting all the 3 numeric features as slope is good

```
In [ ]: plt.figure(figsize = (20,5))
num_cols = ['MonthlyServiceCharges', 'TotalAmount', 'tenure']
for col in enumerate(num_cols):
    plt.subplot(1,3,col[0]+1)
    sns.histplot(data = x_num, x = col[1], kde = True, hue = y['Churn'])
    plt.xticks(rotation = 45)
```



scaling the numeric fields

```
In [ ]: #x_num_scaled = x_num[['MonthlyServiceCharges', 'TotalAmount']]
        scaler = MinMaxScaler()
        #x_num = x_num.values.reshape(1, -1)
        x_num[['MonthlyServiceCharges']] = scaler.fit_transform(x_num[['MonthlyServiceCharges']])
        x_num[['TotalAmount']] = scaler.fit_transform(x_num[['TotalAmount']])
        x_num.head()
```

```
Out[ ]:   SeniorCitizen  tenure  MonthlyServiceCharges  TotalAmount
0           0.0      0.0           0.013988      0.027538
1           0.0      5.0           0.366917      0.458611
2           0.0      4.0           0.000000      0.133713
3           0.0      5.0           0.075320      0.215031
4           0.0      0.0           0.602024      0.111910
```

data visulation of categorical features

```
In [ ]: x_cat.info()

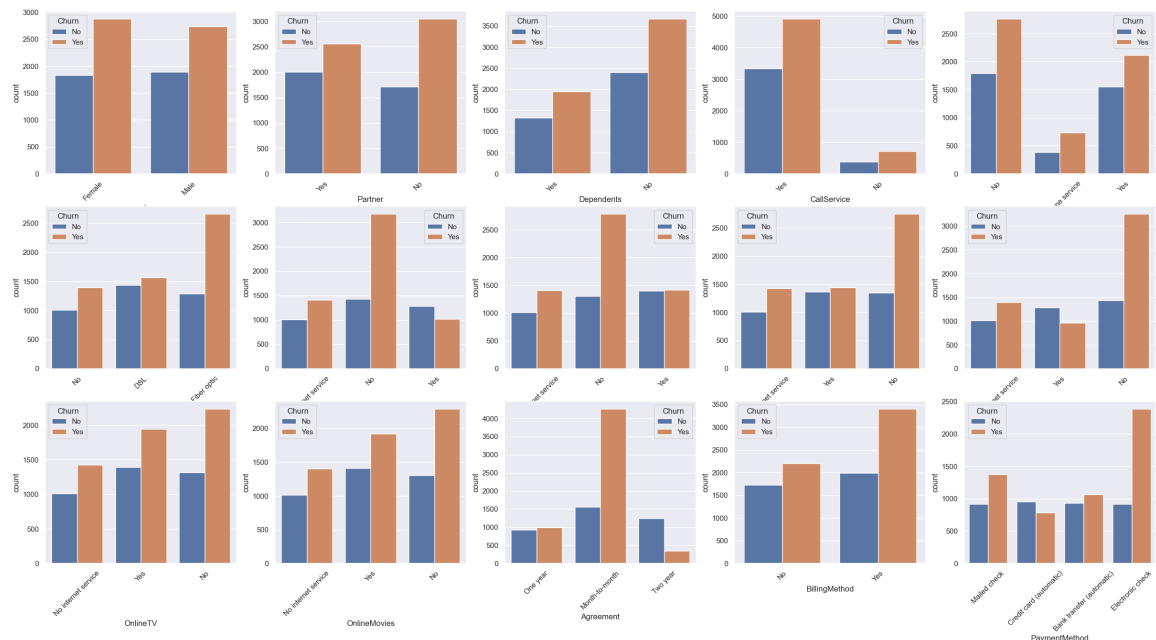
<class 'pandas.core.frame.DataFrame'>
Int64Index: 9328 entries, 0 to 12334
Data columns (total 15 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   gender                                9328 non-null   object
 1   Partner                               9328 non-null   object
 2   Dependents                           9328 non-null   object
 3   CallService                          9328 non-null   object
 4   MultipleConnections                  9328 non-null   object
 5   InternetConnection                   9328 non-null   object
 6   OnlineSecurity                       9328 non-null   object
 7   OnlineBackup                         9328 non-null   object
 8   DeviceProtectionService              9328 non-null   object
 9   TechnicalHelp                        9328 non-null   object
10   OnlineTV                             9328 non-null   object
11   OnlineMovies                         9328 non-null   object
12   Agreement                            9328 non-null   object
13   BillingMethod                        9328 non-null   object
14   PaymentMethod                       9328 non-null   object
dtypes: object(15)
memory usage: 1.1+ MB
```

```
In [ ]: cat_cols = x_cat.columns
        list(enumerate(cat_cols))
```



```
Out[ ]: [(0, 'gender'),
(1, 'Partner'),
(2, 'Dependents'),
(3, 'CallService'),
(4, 'MultipleConnections'),
(5, 'InternetConnection'),
(6, 'OnlineSecurity'),
(7, 'OnlineBackup'),
(8, 'DeviceProtectionService'),
(9, 'TechnicalHelp'),
(10, 'OnlineTV'),
(11, 'OnlineMovies'),
(12, 'Agreement'),
(13, 'BillingMethod'),
(14, 'PaymentMethod')]
```

```
In [ ]: plt.figure(figsize = (30,15))
for col in enumerate(cat_cols):
    plt.subplot(3,5,col[0]+1)
    sns.countplot(x=col[1], hue=y['Churn'], data=x_cat)
    plt.xticks(rotation =45)
```



dropping gender, partner as there is any big variance

```
In [ ]: x_cat = x_cat.drop(['gender', 'Partner'], axis=1)
x_cat.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 9328 entries, 0 to 12334
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Dependents                            9328 non-null   object
1   CallService                           9328 non-null   object
2   MultipleConnections                    9328 non-null   object
3   InternetConnection                    9328 non-null   object
4   OnlineSecurity                         9328 non-null   object
5   OnlineBackup                          9328 non-null   object
6   DeviceProtectionService                9328 non-null   object
7   TechnicalHelp                          9328 non-null   object
8   OnlineTV                              9328 non-null   object
9   OnlineMovies                          9328 non-null   object
10  Agreement                             9328 non-null   object
11  BillingMethod                          9328 non-null   object
12  PaymentMethod                         9328 non-null   object
dtypes: object(13)
memory usage: 1020.2+ KB
```

encoding categorical - using label encoding

```
In [ ]: x_cat_enc = pd.get_dummies(x_cat, drop_first=True)
x_cat_enc.head()
```

```
Out[ ]:
```

	Dependents_Yes	CallService_Yes	MultipleConnections_No phone service	MultipleConnections_Yes	InternetConnection_Yes
0	1	1	0	0	
1	0	0	1	0	
2	0	1	0	0	
3	1	1	0	1	
4	0	1	0	1	

5 rows × 24 columns

selecting k-best features

```
In [ ]: # Select K Best for Categorical Features

selector = SelectKBest(chi2, k=24)
selector.fit_transform(x_cat_enc, y)
# Get columns to keep and create new dataframe with those only
cols = selector.get_support(indices=True)
select_features_df_char = x_cat_enc.iloc[:, cols]
select_features_df_char
```

Out[]:

	Dependents_Yes	CallService_Yes	MultipleConnections_No phone service	MultipleConnections_Yes	Inte
0	1	1	0	0	
1	0	0	1	0	
2	0	1	0	0	
3	1	1	0	1	
4	0	1	0	1	
...
12330	1	1	0	0	
12331	0	0	1	0	
12332	1	0	1	0	
12333	1	1	0	0	
12334	0	1	0	0	

9328 rows × 24 columns

In []:

select_features_df_char.info()

<class 'pandas.core.frame.DataFrame'>

Int64Index: 9328 entries, 0 to 12334

Data columns (total 24 columns):

#	Column	Non-Null Count	Dtype
0	Dependents_Yes	9328 non-null	uint8
1	CallService_Yes	9328 non-null	uint8
2	MultipleConnections_No phone service	9328 non-null	uint8
3	MultipleConnections_Yes	9328 non-null	uint8
4	InternetConnection_Fiber optic	9328 non-null	uint8
5	InternetConnection_No	9328 non-null	uint8
6	OnlineSecurity_No internet service	9328 non-null	uint8
7	OnlineSecurity_Yes	9328 non-null	uint8
8	OnlineBackup_No internet service	9328 non-null	uint8
9	OnlineBackup_Yes	9328 non-null	uint8
10	DeviceProtectionService_No internet service	9328 non-null	uint8
11	DeviceProtectionService_Yes	9328 non-null	uint8
12	TechnicalHelp_No internet service	9328 non-null	uint8
13	TechnicalHelp_Yes	9328 non-null	uint8
14	OnlineTV_No internet service	9328 non-null	uint8
15	OnlineTV_Yes	9328 non-null	uint8
16	OnlineMovies_No internet service	9328 non-null	uint8
17	OnlineMovies_Yes	9328 non-null	uint8
18	Agreement_One year	9328 non-null	uint8
19	Agreement_Two year	9328 non-null	uint8
20	BillingMethod_Yes	9328 non-null	uint8
21	PaymentMethod_Credit card (automatic)	9328 non-null	uint8
22	PaymentMethod_Electronic check	9328 non-null	uint8
23	PaymentMethod_Mailed check	9328 non-null	uint8

dtypes: uint8(24)

memory usage: 291.5 KB

final features

```
In [ ]: x_cols_final = pd.concat([x_num,select_features_df_char],axis=1, join="inner")
x_cols_final.columns
```

```
Out[ ]: Index(['SeniorCitizen', 'tenure', 'MonthlyServiceCharges', 'TotalAmount',
              'Dependents_Yes', 'CallService_Yes',
              'MultipleConnections_No phone service', 'MultipleConnections_Yes',
              'InternetConnection_Fiber optic', 'InternetConnection_No',
              'OnlineSecurity_No internet service', 'OnlineSecurity_Yes',
              'OnlineBackup_No internet service', 'OnlineBackup_Yes',
              'DeviceProtectionService_No internet service',
              'DeviceProtectionService_Yes', 'TechnicalHelp_No internet service',
              'TechnicalHelp_Yes', 'OnlineTV_No internet service', 'OnlineTV_Yes',
              'OnlineMovies_No internet service', 'OnlineMovies_Yes',
              'Agreement_One year', 'Agreement_Two year', 'BillingMethod_Yes',
              'PaymentMethod_Credit card (automatic)',
              'PaymentMethod_Electronic check', 'PaymentMethod_Mailed check'],
              dtype='object')
```

```
In [ ]: x_cols_final.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 9328 entries, 0 to 12334
Data columns (total 28 columns):
 #   Column                                                                 Non-Null Count  Dtype
---  -
 0   SeniorCitizen                                                         9328 non-null   float64
 1   tenure                                                                9328 non-null   float64
 2   MonthlyServiceCharges                                                 9328 non-null   float64
 3   TotalAmount                                                           9328 non-null   float64
 4   Dependents_Yes                                                        9328 non-null   uint8
 5   CallService_Yes                                                       9328 non-null   uint8
 6   MultipleConnections_No phone service                                9328 non-null   uint8
 7   MultipleConnections_Yes                                               9328 non-null   uint8
 8   InternetConnection_Fiber optic                                        9328 non-null   uint8
 9   InternetConnection_No                                                9328 non-null   uint8
10   OnlineSecurity_No internet service                                    9328 non-null   uint8
11   OnlineSecurity_Yes                                                    9328 non-null   uint8
12   OnlineBackup_No internet service                                       9328 non-null   uint8
13   OnlineBackup_Yes                                                      9328 non-null   uint8
14   DeviceProtectionService_No internet service                         9328 non-null   uint8
15   DeviceProtectionService_Yes                                           9328 non-null   uint8
16   TechnicalHelp_No internet service                                     9328 non-null   uint8
17   TechnicalHelp_Yes                                                     9328 non-null   uint8
18   OnlineTV_No internet service                                          9328 non-null   uint8
19   OnlineTV_Yes                                                          9328 non-null   uint8
20   OnlineMovies_No internet service                                       9328 non-null   uint8
21   OnlineMovies_Yes                                                      9328 non-null   uint8
22   Agreement_One year                                                    9328 non-null   uint8
23   Agreement_Two year                                                    9328 non-null   uint8
24   BillingMethod_Yes                                                      9328 non-null   uint8
25   PaymentMethod_Credit card (automatic)                                9328 non-null   uint8
26   PaymentMethod_Electronic check                                        9328 non-null   uint8
27   PaymentMethod_Mailed check                                            9328 non-null   uint8
dtypes: float64(4), uint8(24)
memory usage: 841.0 KB
```

```
In [ ]: x_cols_final.nunique()
```

```
Out[ ]: SeniorCitizen      2
        tenure            6
        MonthlyServiceCharges  5153
        TotalAmount        8831
        Dependents_Yes      2
        CallService_Yes     2
        MultipleConnections_No phone service  2
        MultipleConnections_Yes  2
        InternetConnection_Fiber optic  2
        InternetConnection_No  2
        OnlineSecurity_No internet service  2
        OnlineSecurity_Yes     2
        OnlineBackup_No internet service  2
        OnlineBackup_Yes      2
        DeviceProtectionService_No internet service  2
        DeviceProtectionService_Yes  2
        TechnicalHelp_No internet service  2
        TechnicalHelp_Yes     2
        OnlineTV_No internet service  2
        OnlineTV_Yes          2
        OnlineMovies_No internet service  2
        OnlineMovies_Yes      2
        Agreement_One year    2
        Agreement_Two year    2
        BillingMethod_Yes      2
        PaymentMethod_Credit card (automatic)  2
        PaymentMethod_Electronic check  2
        PaymentMethod_Mailed check  2
        dtype: int64
```

```
In [ ]: y['Churn'] = np.where(y['Churn']=='No',0,1)
        y['Churn'].head()
```

```
Out[ ]: 0    0
        1    0
        2    0
        3    0
        4    0
        Name: Churn, dtype: int32
```

variance influence factor for x_cols_final.
we observe that there exists very high correlation between tenure, MonthlyServiceCharges, TotalAmount, CallService_Yes features having very high correlation. Also not eliminating any feature here

```
In [ ]: vif_features = pd.DataFrame()
        vif_features['feature_names'] = x_cols_final.columns
```

```
vif_features['vif_values'] = [variance_inflation_factor(x_cols_final, i) for i in range(x_cols_final.shape[0])]
print(vif_features)
```

	feature_names	vif_values
0	SeniorCitizen	1.283691
1	tenure	11.164181
2	MonthlyServiceCharges	17.803474
3	TotalAmount	15.518206
4	Dependents_Yes	1.621904
5	CallService_Yes	11.620357
6	MultipleConnections_No phone service	1.589235
7	MultipleConnections_Yes	2.196782
8	InternetConnection_Fiber optic	4.006358
9	InternetConnection_No	3.949974
10	OnlineSecurity_No internet service	3.369774
11	OnlineSecurity_Yes	1.781707
12	OnlineBackup_No internet service	3.389887
13	OnlineBackup_Yes	1.965955
14	DeviceProtectionService_No internet service	3.458866
15	DeviceProtectionService_Yes	2.039356
16	TechnicalHelp_No internet service	3.284803
17	TechnicalHelp_Yes	1.855395
18	OnlineTV_No internet service	3.566204
19	OnlineTV_Yes	2.504054
20	OnlineMovies_No internet service	3.535424
21	OnlineMovies_Yes	2.458684
22	Agreement_One year	1.599268
23	Agreement_Two year	2.228679
24	BillingMethod_Yes	2.514791
25	PaymentMethod_Credit card (automatic)	1.772461
26	PaymentMethod_Electronic check	2.628250
27	PaymentMethod_Mailed check	2.104484

splitting train and test

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(x_cols_final, y, test_size=0.2, random_state=42)
model_accuracy = {}
#model_accuracy = {'LogisticRegression':[0], 'RandomForestClassifier':[0], 'GridSearchCV':[0]}
```

```
In [ ]: print("Shape of Training Data",X_train.shape)
print("Shape of Testing Data",X_test.shape)
print("Response Rate in Training Data",y_train.mean())
print("Response Rate in Testing Data",y_test.mean())

# observe that the mean of training and test of y_train and y_test are more or less same
```

```
Shape of Training Data (6529, 28)
Shape of Testing Data (2799, 28)
Response Rate in Training Data Churn    0.604074
dtype: float64
Response Rate in Testing Data Churn    0.594141
dtype: float64
```

model execution and summary with coefficients & p-values

```
In [ ]: import statsmodels.api as sm
        #y_train = list(y_train)
        model_result = sm.Logit(y_train,X_train).fit()
        print(model_result.summary())
```

Optimization terminated successfully.

Current function value: 0.538566

Iterations 6

Logit Regression Results

```
=====
Dep. Variable:          Churn    No. Observations:          6529
Model:                  Logit    Df Residuals:              6501
Method:                  MLE     Df Model:                  27
Date:                    Sun, 26 Mar 2023    Pseudo R-squ.:          0.1978
Time:                    13:34:48    Log-Likelihood:         -3516.3
converged:                True    LL-Null:                 -4383.1
Covariance Type:          nonrobust    LLR p-value:            0.000
=====
```

```
=====
                                coef    std err          z
P>|z|      [0.025    0.975]
-----
SeniorCitizen          -0.2541    0.095    -2.681
0.007    -0.440    -0.068
tenure          -0.3301    0.040    -8.289
0.000    -0.408    -0.252
MonthlyServiceCharges    1.4090    0.234    6.010
0.000    0.950    1.868
TotalAmount          -1.0628    0.345    -3.084
0.002    -1.738    -0.387
Dependents_Yes          0.3082    0.063    4.864
0.000    0.184    0.432
CallService_Yes          0.0511    0.103    0.498
0.619    -0.150    0.252
MultipleConnections_No phone service    0.9166    0.108    8.500
0.000    0.705    1.128
MultipleConnections_Yes    0.2303    0.070    3.281
0.001    0.093    0.368
InternetConnection_Fiber optic    0.5991    0.088    6.770
0.000    0.426    0.773
InternetConnection_No    0.3432    0.111    3.090
0.002    0.125    0.561
OnlineSecurity_No internet service    0.0224    0.103    0.217
0.828    -0.180    0.225
OnlineSecurity_Yes          -0.3181    0.076    -4.213
0.000    -0.466    -0.170
OnlineBackup_No internet service    0.0753    0.103    0.731
0.465    -0.127    0.277
OnlineBackup_Yes          -0.0783    0.074    -1.058
0.290    -0.223    0.067
DeviceProtectionService_No internet service    0.0806    0.104    0.771
0.440    -0.124    0.285
DeviceProtectionService_Yes    0.0478    0.076    0.629
0.529    -0.101    0.197
TechnicalHelp_No internet service    -0.0530    0.102    -0.520
0.603    -0.253    0.147
TechnicalHelp_Yes          -0.3501    0.079    -4.457
0.000    -0.504    -0.196
OnlineTV_No internet service    0.1919    0.105    1.836
0.066    -0.013    0.397
OnlineTV_Yes          0.1198    0.078    1.532
0.126    -0.033    0.273
OnlineMovies_No internet service    0.1157    0.106    1.094
0.274    -0.092    0.323
```


OnlineMovies_Yes	0.1107	0.077	1.430
0.153 -0.041 0.263			
Agreement_One year	-0.2188	0.076	-2.881
0.004 -0.368 -0.070			
Agreement_Two year	-1.1217	0.104	-10.834
0.000 -1.325 -0.919			
BillingMethod_Yes	0.1704	0.060	2.832
0.005 0.052 0.288			
PaymentMethod_Credit card (automatic)	-0.0805	0.088	-0.914
0.361 -0.253 0.092			
PaymentMethod_Electronic check	0.3306	0.079	4.202
0.000 0.176 0.485			
PaymentMethod_Mailed check	-0.0581	0.083	-0.704
0.481 -0.220 0.104			

=====

=====

creating the model

```
In [ ]: logreg = LogisticRegression()
logreg.fit(X_train, y_train)
```

```
Out[ ]: ▾ LogisticRegression
LogisticRegression()
```


predicting the test set results and calculating the accuracy

```
In [ ]: y_pred_test = logreg.predict(X_test)
score = logreg.score(X_test, y_test)
print('Accuracy score:', score)
model_accuracy['LogisticRegression'] = score
```

Accuracy score: 0.7227581279028225

```
In [ ]: confusion_matrix = metrics.confusion_matrix(y_test, y_pred_test)
plt.figure(figsize = (3,2))
sns.heatmap(confusion_matrix , annot=True,fmt = "d", cmap='OrRd')
plt.title("LOGISTIC REGRESSION CONFUSION MATRIX");
```

LOGISTIC REGRESSION CONFUSION MATRIX



classification report

```
In [ ]: print("classification_report")
print(metrics.classification_report(y_test, y_pred_test))
```

```
classification_report
              precision    recall  f1-score   support

     0       0.72       0.53       0.61       1136
     1       0.73       0.86       0.79       1663

 accuracy                   0.72       2799
 macro avg       0.72       0.69       0.70       2799
weighted avg       0.72       0.72       0.71       2799
```

Training Accuracy vs Test Accuracy, both seemed to be more or less same in accuracy

```
In [ ]: y_pred_train = logreg.predict(X_train)
print("Training Accuracy", metrics.accuracy_score(y_train, y_pred_train))
print("*****"*5)
print("Test Accuracy", metrics.accuracy_score(y_test, y_pred_test))
```

```
Training Accuracy 0.739163731046102
*****
*
Test Accuracy 0.7227581279028225
```


ROC - AUC

```
In [ ]: logistic_roc_auc = metrics.roc_auc_score(y_test, y_pred_test)
logistic_roc_auc
```

```
Out[ ]: 0.6916467990141693
```

```
In [ ]: fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_test)
```

```
In [ ]: plt.figure()
plt.plot(fpr, tpr, label = "ROC Curve (area=%0.2f)" %logistic_roc_auc)
plt.plot([0,1],[0,1], 'k--')
plt.xlim([0.0,1.0])
plt.ylim([0.0,1.0])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Receiver Operating Characteristic with Area Under Curve")
plt.legend(loc='lower right')
plt.show()
```



performing cross validation and observed that accuracy max is 80% and min 68%

```
In [ ]: # cross validation score - K_FOLD method
from sklearn.model_selection import cross_val_score
accuracy = cross_val_score(logreg, X_test, y_test, cv =20)
print(accuracy)
print('accuracy mean ', accuracy.mean())

[0.72142857 0.72857143 0.72857143 0.75714286 0.74285714 0.75714286
 0.72142857 0.69285714 0.77857143 0.66428571 0.76428571 0.77142857
 0.78571429 0.7          0.69285714 0.73571429 0.73571429 0.72857143
 0.73571429 0.6618705 ]
accuracy mean  0.7302363823227133
```

using Random forest and found that accuracy improved to 78%


```
In [ ]: rclassifier = RandomForestClassifier(n_estimators=25, criterion='gini', random_s
rclassifier.fit(X_train, y_train)
```

```
Out[ ]: RandomForestClassifier
RandomForestClassifier(max_depth=5, n_estimators=25, random_state=1)
```

```
In [ ]: y_pred_test_rfc = rclassifier.predict(X_test)
score_rfc = metrics.accuracy_score(y_test, y_pred_test_rfc)
model_accuracy['RandomForestClassifier'] = score_rfc
```

```
In [ ]: confusion_matrix = metrics.confusion_matrix(y_test, y_pred_test_rfc)
plt.figure(figsize = (3,2))
sns.heatmap(confusion_matrix , annot=True,fmt = "d", cmap='OrRd')
plt.title("RANDOM FOREST CONFUSION MATRIX")
```

```
Out[ ]: Text(0.5, 1.0, 'RANDOM FOREST CONFUSION MATRIX')
```



print classification report

```
In [ ]: print("classification_report")
print(metrics.classification_report(y_test, y_pred_test_rfc))
```

```
classification_report
              precision    recall  f1-score   support

     0       0.82         0.60         0.70         1136
     1       0.77         0.91         0.83         1663

 accuracy          0.79
 macro avg         0.80         0.76         0.77         2799
weighted avg         0.79         0.79         0.78         2799
```

RandomForestClassifier cross validation

```
In [ ]: accuracy = cross_val_score(rclassifier, X_test, y_test, cv =20)
print(accuracy)
print('accuracy mean ', accuracy.mean())

[0.80714286 0.81428571 0.74285714 0.76428571 0.77142857 0.78571429
 0.79285714 0.76428571 0.83571429 0.72142857 0.77857143 0.83571429
 0.78571429 0.77857143 0.77857143 0.8          0.76428571 0.79285714
 0.79285714 0.79856115]
accuracy mean  0.7852852004110996
```

Random forest feature importance

```
In [ ]: features_imp_dict= {}
for col, val in sorted(zip(x_cols_final.columns, rclassifier.feature_importances_)):
    features_imp_dict[col]=val
features_df_rf = pd.DataFrame({'Feature':features_imp_dict.keys(),'Importance':features_imp_dict.values()})
features_df_rf
```


Out[]:

	Feature	Importance
0	tenure	0.211375
1	MonthlyServiceCharges	0.187146
2	Agreement_Two year	0.145862
3	TotalAmount	0.128925
4	OnlineMovies_No internet service	0.038637
5	InternetConnection_Fiber optic	0.037694
6	OnlineSecurity_Yes	0.035808
7	InternetConnection_No	0.028067
8	TechnicalHelp_No internet service	0.027064
9	DeviceProtectionService_No internet service	0.024846
10	OnlineSecurity_No internet service	0.022896
11	OnlineBackup_No internet service	0.021353
12	PaymentMethod_Electronic check	0.019439
13	TechnicalHelp_Yes	0.015250
14	OnlineTV_No internet service	0.015148
15	OnlineBackup_Yes	0.011226
16	PaymentMethod_Credit card (automatic)	0.006525
17	MultipleConnections_Yes	0.004846
18	Agreement_One year	0.004833
19	DeviceProtectionService_Yes	0.002495
20	MultipleConnections_No phone service	0.002212
21	OnlineTV_Yes	0.002154
22	BillingMethod_Yes	0.002125
23	OnlineMovies_Yes	0.001145
24	Dependents_Yes	0.000873
25	PaymentMethod_Mailed check	0.000835
26	CallService_Yes	0.000758
27	SeniorCitizen	0.000464

	Feature	Importance
0	tenure	0.211375
1	MonthlyServiceCharges	0.187146
2	Agreement_Two year	0.145862
3	TotalAmount	0.128925
4	OnlineMovies_No internet service	0.038637
5	InternetConnection_Fiber optic	0.037694
6	OnlineSecurity_Yes	0.035808
7	InternetConnection_No	0.028067
8	TechnicalHelp_No internet service	0.027064
9	DeviceProtectionService_No internet service	0.024846
10	OnlineSecurity_No internet service	0.022896
11	OnlineBackup_No internet service	0.021353
12	PaymentMethod_Electronic check	0.019439
13	TechnicalHelp_Yes	0.015250
14	OnlineTV_No internet service	0.015148
15	OnlineBackup_Yes	0.011226
16	PaymentMethod_Credit card (automatic)	0.006525
17	MultipleConnections_Yes	0.004846
18	Agreement_One year	0.004833
19	DeviceProtectionService_Yes	0.002495
20	MultipleConnections_No phone service	0.002212
21	OnlineTV_Yes	0.002154
22	BillingMethod_Yes	0.002125
23	OnlineMovies_Yes	0.001145
24	Dependents_Yes	0.000873
25	PaymentMethod_Mailed check	0.000835
26	CallService_Yes	0.000758
27	SeniorCitizen	0.000464

```
In [ ]: values = features_df_rf.Importance
idx = features_df_rf.Feature
plt.figure(figsize=(10,8))
clrs = ['green' if (x < max(values)) else 'red' for x in values ]
```

```
sns.barplot(y=idx,x=values,palette=clrs).set(title='Important features to predic
plt.show()
```



using GridSearchCV and Random Forest

```
In [ ]: # selected model up on accuracy score in RandomForestClassifier with Tunned Hyper
params = {'n_estimators' : [30],
          'min_samples_split' : [8, 10, 12],
          'criterion' : ['gini', 'entropy']}

rfc_gscv = GridSearchCV(RandomForestClassifier(),param_grid= params, scoring='roc_auc')
rfc_gscv.fit(X_train, y_train)
y_predicted = rfc_gscv.predict(X_test)

rfc_gscv_model = rfc_gscv.best_estimator_
print (rfc_gscv.best_score_, rfc_gscv.best_params_)
```

```
0.9113187256349988 {'criterion': 'gini', 'min_samples_split': 10, 'n_estimators': 30}
```

```
In [ ]: y_pred_gscv = rfc_gscv.predict(X_test)
score_gscv = accuracy_score(y_test, y_pred_gscv)

model_accuracy['GridSearchCV_RFC'] = score_gscv

print("Grid Search CV", score_gscv)
print("classification_report")
print(metrics.classification_report(y_test, y_pred_gscv))
```

Grid Search CV 0.8113612004287245


classification_report

	precision	recall	f1-score	support
0	0.78	0.75	0.76	1136
1	0.83	0.85	0.84	1663
accuracy			0.81	2799
macro avg	0.81	0.80	0.80	2799
weighted avg	0.81	0.81	0.81	2799

```
In [ ]: confusion_matrix = metrics.confusion_matrix(y_test, y_pred_gscv)
plt.figure(figsize = (3,2))
sns.heatmap(confusion_matrix , annot=True,fmt = "d", cmap='OrRd')
plt.title("GRIDSEARCHCV RANDOM CLASSIFIER CONFUSION MATRIX")
```

```
Out[ ]: Text(0.5, 1.0, 'GRIDSEARCHCV RANDOM CLASSIFIER CONFUSION MATRIX')
```

GRIDSEARCHCV RANDOM CLASSIFIER CONFUSION MATRIX



using gradient boosting classifier (reducing the errors) and further improvement on accuracy to 82%

```
In [ ]: gbc_classifier = GradientBoostingClassifier(random_state=0,max_depth=5,min_sampl
gbc_classifier.fit(X_train, y_train)
y_pred_gbc1 = gbc_classifier.predict(X_test)

score_gbc1 = accuracy_score(y_test, y_pred_gbc1)
model_accuracy['GradientBoostingClassifier'] = score_gbc1

print("Gradient Boosting Classifier", accuracy_score(y_test, y_pred_gbc1))
```

Gradient Boosting Classifier 0.8177920685959271

```
In [ ]: y_pred_gbc1 = gbc_classifier.predict(X_test)
print("Gradient Boosting Classifier", accuracy_score(y_test, y_pred_gbc1))
print("classification_report")
print(metrics.classification_report(y_test, y_pred_gbc1))
```

Gradient Boosting Classifier 0.8177920685959271


classification_report

	precision	recall	f1-score	support
0	0.78	0.76	0.77	1136
1	0.84	0.86	0.85	1663
accuracy			0.82	2799
macro avg	0.81	0.81	0.81	2799
weighted avg	0.82	0.82	0.82	2799

```
In [ ]: confusion_matrix = metrics.confusion_matrix(y_test, y_pred_gbc1)
plt.figure(figsize = (3,2))
sns.heatmap(confusion_matrix , annot=True,fmt = "d", cmap='OrRd')
plt.title("GRADIANT BOOSTING CLASSIFIER CONFUSION MATRIX")
```

```
Out[ ]: Text(0.5, 1.0, 'GRADIANT BOOSTING CLASSIFIER CONFUSION MATRIX')
```

GRADIANT BOOSTING CLASSIFIER CONFUSION MATRIX



understanding the feature importance using gbc

```
In [ ]: # Let's create a dictionary of features and their importance values
features_imp_dict= {}
for col, val in sorted(zip(x_cols_final.columns, gbc_classifier.feature_importances_)):
    features_imp_dict[col]=val
features_df = pd.DataFrame({'Feature':features_imp_dict.keys(),'Importance':features_imp_dict.values()})
features_df
```


Out[]:

	Feature	Importance
0	MonthlyServiceCharges	0.266159
1	tenure	0.208398
2	TotalAmount	0.095333
3	Agreement_Two year	0.090201
4	InternetConnection_No	0.032979
5	OnlineTV_No internet service	0.031074
6	InternetConnection_Fiber optic	0.030000
7	OnlineBackup_No internet service	0.029369
8	OnlineMovies_No internet service	0.023188
9	TechnicalHelp_No internet service	0.022890
10	DeviceProtectionService_No internet service	0.022751
11	MultipleConnections_No phone service	0.021824
12	OnlineSecurity_No internet service	0.018124
13	TechnicalHelp_Yes	0.017099
14	Agreement_One year	0.016186
15	PaymentMethod_Electronic check	0.012911
16	OnlineSecurity_Yes	0.011377
17	CallService_Yes	0.009090
18	MultipleConnections_Yes	0.009051
19	OnlineMovies_Yes	0.007741
20	OnlineTV_Yes	0.007339
21	SeniorCitizen	0.004076
22	OnlineBackup_Yes	0.003640
23	DeviceProtectionService_Yes	0.003506
24	BillingMethod_Yes	0.002522
25	PaymentMethod_Credit card (automatic)	0.001586
26	Dependents_Yes	0.001127
27	PaymentMethod_Mailed check	0.000457


visualizing the important features, gbc

In []:

```

values = features_df.Importance
idx = features_df.Feature
plt.figure(figsize=(10,8))
clrs = ['green' if (x < max(values)) else 'red' for x in values ]
sns.barplot(y=idx,x=values,palette=clrs).set(title='Important features to predic
plt.show()

```



using RidgeClassifier to observe the model

```
In [ ]: rgclassifier = RidgeClassifier()
        rgclassifier.fit(X_train, y_train)
```

```
Out[ ]: ▾ RidgeClassifier
        RidgeClassifier()
```

```
In [ ]: rgclassifier.fit(X_train, y_train)
        y_pred_train_rdg = rgclassifier.predict(X_train)
        score_train_rdg = accuracy_score(y_train, y_pred_train_rdg)
        print('RidgeClassifier training accuracy', score_train_rdg)

        y_pred_test_rdg = rgclassifier.predict(X_test)
        score_test_rdg = accuracy_score(y_test, y_pred_test_rdg)
        model_accuracy['RidgeClassifier'] = score_test_rdg
        print('RidgeClassifier testing accuracy', score_test_rdg)
```

```
RidgeClassifier training accuracy 0.7377852657374789
RidgeClassifier testing accuracy 0.7220435869953555
```

```
In [ ]: print("classification_report")
        print(metrics.classification_report(y_test, y_pred_test_rdg))
```

```

classification_report
              precision    recall  f1-score   support

     0       0.72       0.52       0.60       1136
     1       0.72       0.86       0.79       1663

 accuracy          0.72          2799
 macro avg       0.72       0.69       0.69       2799
 weighted avg    0.72       0.72       0.71       2799

```

```


In [ ]: confusion_matrix = metrics.confusion_matrix(y_test, y_pred_test_rdg)
plt.figure(figsize = (3,2))
sns.heatmap(confusion_matrix , annot=True,fmt = "d", cmap='OrRd')
plt.title("RidgeClassifier CONFUSION MATRIX")

```

```

Out[ ]: Text(0.5, 1.0, 'RidgeClassifier CONFUSION MATRIX')

```



using SVC

```

In [ ]: svc = SVC()
svc.fit(X_train, y_train)
score = svc.score(X_train, y_train)
print("SVC training dataset: ", score)

cv_scores = cross_val_score(svc, X_test, y_test, cv=10)
print("CV average score: %.2f" % cv_scores.mean())

y_pred_svc = svc.predict(X_test)
score_test_svc = accuracy_score(y_test, y_pred_svc)
model_accuracy['SVC'] = score_test_svc
print('SVC testing accuracy', score_test_svc)

confusion_matrix = metrics.confusion_matrix(y_test, y_pred_svc)
plt.figure(figsize = (3,2))
sns.heatmap(confusion_matrix , annot=True,fmt = "d", cmap='OrRd')
plt.title("SVC CONFUSION MATRIX")

print("SVC classification_report")
print(metrics.classification_report(y_test, y_pred_svc))

```


SVC training dataset: 0.8286108132945321

CV average score: 0.80

SVC testing accuracy 0.8149339049660593

SVC classification_report

	precision	recall	f1-score	support
0	0.81	0.72	0.76	1136
1	0.82	0.88	0.85	1663
accuracy			0.81	2799
macro avg	0.81	0.80	0.80	2799
weighted avg	0.81	0.81	0.81	2799



GridSearchCV(RandomForest), GradientBoostingClassifier, SVC has better accuracy

In []: model_accuracy

Out[]: {'LogisticRegression': 0.7227581279028225,
 'RandomForestClassifier': 0.7859949982136477,
 'GridSearchCV_RFC': 0.8113612004287245,
 'GradientBoostingClassifier': 0.8177920685959271,
 'RidgeClassifier': 0.7220435869953555,
 'SVC': 0.8149339049660593}